

Ticket 12 Study Report: Solver Validation Synthesis / Review Gate

SURF2026 American Option Risk Surfaces

Codex-assisted implementation report

June 16, 2026

Abstract

Ticket 12 synthesizes the solver-validation evidence built across Tickets 01 through 11. It does not add a new pricing method. It audits existing artifacts, summarizes the validation ladder, and makes a conservative review-gate decision for the classical American Crank-Nicolson / projected SOR solver. The decision is `PASS_SOLVER_VALIDATION_WITH_LIMITATIONS`. This means the current one-dimensional Black-Scholes American CN/PSOR baseline is credible enough for formal reference integration and human review, while stress maps, datasets, neural surrogate labels, production Greeks, and final bibliography integration remain out of scope.

Contents

1	Role of Ticket 12 in the Whole Project	3
2	Validation Ladder Overview	3
3	Methodology and Source Basis	4
4	Summary of Tickets 01–11	5
5	Acceptance Criteria Review	5
6	Evidence Summary	6
7	Price Validation Evidence	7
8	American Obstacle and LCP Evidence	7
9	Boundary Extraction Evidence	7
10	Greek Diagnostic Evidence	8
11	Rannacher Smoothing Gate Evidence	8

12 Grid and Domain Sensitivity Evidence	8
13 Artifact Audit Summary	8
14 Gate Decision	9
15 Claims Supported by Current Evidence	9
16 Claims Not Yet Supported	10
17 Implementation Pitfalls and Debugging Notes	10
18 Limitations	11
19 Lessons Learned / Study Notes	11
20 Recommended Next Steps	12
21 Reproducibility Record	12

1 Role of Ticket 12 in the Whole Project

The project is building toward American option risk surfaces: prices, early-exercise boundaries, Delta, Gamma, and later stress or surrogate-model studies. The planning report and solver validation plan both make one point repeatedly: downstream work should wait until the classical solver has been checked carefully. If the solver is wrong, then any later stress map or machine-learning model can only learn unreliable labels.

Ticket 12 is the review gate after the first solver-validation ladder. It asks:

Is the classical American CN/PSOR solver validated enough for the next research stage, and what limitations must be kept?

The answer in this report is yes, but with important limits. The current evidence supports moving to reference integration and human review. It does not support jumping straight to large-scale stress maps, dataset construction, neural-network labels, or production risk reporting.

2 Validation Ladder Overview

The validation ladder was intentionally incremental. Early tickets checked analytic utilities and finite-difference building blocks. Middle tickets validated European and American option behavior. Later tickets added diagnostics, boundary extraction, Greeks, Rannacher comparison, and grid/domain sensitivity.

Ticket	Purpose	Main evidence	Status	Main limitation
01	Payoffs and European Black-Scholes utilities	Payoff, parity, scalar/vector, $T = 0$, and $\sigma = 0$ tests.	PASS	Analytic utilities only.
02	Uniform grids and Black-Scholes operator	Grid endpoints, operator coefficients, boundary helpers, manufactured-function checks.	PASS	Uniform central-difference setup only.
03	European CN validation	Closed-form comparison: max target-region error $2.68274272102 \times 10^{-4}$.	PASS	No American obstacle.
04	PSOR/LCP core	Projected SOR core, obstacle smoke tests, convergence metadata.	PASS with limits	Not a full validation study.
05	American put validation	Obstacle violation 0, positive American-over-European value, PSOR convergence.	PASS	Representative case only.
06	No-dividend American call theorem	Fine-grid American-European max error $1.23566318468 \times 10^{-4}$.	PASS	Approximate finite-difference equality.

Ticket	Purpose	Main evidence	Status	Main limitation
07	Dividend-paying American call	Positive American-minus-European value; fine-grid max 0.0894222439428.	PASS	Parameter-specific dividend-call evidence.
08	LCP diagnostics	Max equation violation $2.14416621448 \times 10^{-9}$, complementarity product $6.6483476629 \times 10^{-11}$.	PASS	Diagnostic tolerances, not proof.
09	Boundary extraction	Put and dividend-call boundaries found at 120 positive- τ rows; no-dividend call control found no boundary.	PASS with limits	Threshold-based boundaries.
10	Delta/Gamma diagnostics	Full-grid max $ \Gamma = 30$, strict-mask max $ \Gamma = 3.18001886582$.	PASS with cautions	Greeks are diagnostic only.
10A	Rannacher comparison gate	Gate recommendation KEEP_BASELINE_FOR_NOW ; max price difference $1.72323866424 \times 10^{-6}$.	PASS with limits	Rannacher not adopted as default.
11	Grid/domain sensitivity	Grid selected-price diff 0.0111111111111; domain selected-price diff 0; LCP diagnostics stable.	PASS with cautions	Three grid levels are not a convergence proof.

3 Methodology and Source Basis

This synthesis uses only internal project materials and existing ticket artifacts:

- reports/00_planning/planning_report.md;
- reports/01_literature/literature_map.md;
- reports/02_math/formulation_note.md;
- reports/03_solver/solver_validation_plan.md;
- reports/03_solver/solver_implementation_tickets.md;
- Ticket 01 through Ticket 11 study reports and result CSVs.

No external bibliography is added here. Formal reference integration is explicitly deferred until after this twelve-report study sequence is complete. That later pass should verify complete bibliographic details and connect the implemented validation ladder to the classical American option, finite-difference, LCP/PSOR, Rannacher smoothing, and scientific machine-learning literature.

4 Summary of Tickets 01–11

Tickets 01 and 02 established the small trusted components: payoff functions, European closed-form prices, grids, spatial operators, and boundary helpers. Ticket 03 then checked the European finite-difference time-stepping path before any American exercise logic was added.

Ticket 04 introduced the American obstacle and PSOR/LCP core. Tickets 05, 06, and 07 checked the three finance behaviors needed for a first American solver: American put exercise value, no-dividend American call equality with the European call, and positive-dividend American call early-exercise value evidence.

Tickets 08 and 09 moved from price checks into structure checks: obstacle/equation/complementarity diagnostics and threshold-based continuation-premium boundary extraction. Ticket 10 added finite-difference Delta and Gamma diagnostics, carefully warning that Gamma is fragile near kinks, maturity, and boundary bands. Ticket 10A compared a named Rannacher-smoothed variant and kept the baseline because smoothing did not improve the headline strict-mask Gamma metrics enough to justify changing the validated default. Ticket 11 checked grid and domain sensitivity for prices, boundaries, LCP diagnostics, and Greek diagnostics.

5 Acceptance Criteria Review

Criterion	Evidence source	Result	Interpretation
Analytic payoff and European utilities are trusted.	Tickets 01 tests and report.	PASS	Later validation can use these utilities as references.
Grid/operator construction is explicit and tested.	Ticket 02 tests and report.	PASS	Interior indexing and boundary helpers are documented.
European CN matches closed-form prices.	Ticket 03 CSV.	PASS	Target-region errors are small for the representative case.
American obstacle is enforced.	Tickets 05, 07, 08, 11 CSVs.	PASS	Max obstacle violation is 0 in representative American cases.
No-dividend call does not invent early exercise value.	Ticket 06 CSV and Ticket 09 control.	PASS	American-European error remains finite-difference small; no boundary is forced.
Dividend-paying call can show early-exercise value.	Ticket 07 CSV.	PASS	Positive American-minus-European value appears for $q = 0.08$.
Complementarity diagnostics are near solver tolerance.	Ticket 08 and 11 CSVs.	PASS	LCP residuals are small in representative cases.

Criterion	Evidence source	Result	Interpretation
Boundary extraction is available with metadata.	Ticket 09 CSV.	PASS with limits	Boundary curves are threshold-based, not exact analytical boundaries.
Greeks are diagnosed cautiously.	Ticket 10 and 11 CSVs.	PASS with cautions	Gamma is informative but not production-ready.
Rannacher is considered without changing default solver.	Ticket 10A CSV.	PASS	Baseline remains default for now.
Grid/domain sensitivity does not contradict baseline use.	Ticket 11 CSVs.	PASS with cautions	Evidence is preliminary, not a formal convergence proof.

6 Evidence Summary

The Ticket 12 script generated `ticket_12_solver_validation_evidence_summary.csv` with 29 evidence rows. The following table condenses the key metrics.

Area	Key metric	Value	Caution
European CN	Max target-region absolute error	$2.68274272102 \times 10^{-4}$	No American obstacle yet.
European CN	Max target-region RMSE	$1.15854293529 \times 10^{-4}$	Representative parameters.
American put	Max obstacle violation	0	Not an external reference-price benchmark.
American put	Max American-European put value	0.0487705754993	Case-specific early-exercise value.
No-dividend call	Fine-grid max American-European error	$1.23566318468 \times 10^{-4}$	Approximate equality, not exact arithmetic equality.
Dividend call	Fine-grid max American-European value	0.0894222439428	Positive-dividend case only.
LCP diagnostics	Max equation violation	$2.14416621448 \times 10^{-9}$	Diagnostic threshold, not proof.
LCP diagnostics	Max complementarity product	$6.6483476629 \times 10^{-11}$	Representative cases.
Boundary extraction	Put and dividend-call found rows	120, 120	Threshold-based extraction.
No-boundary control	No-dividend call control	NO_BOUNDARY_FOUND	Confirms extractor does not force a boundary.
Greeks	Full-grid max $ \Gamma $	30	Kink/boundary/maturity sensitive.
Greeks	Strict-mask max $ \Gamma $	3.18001886582	Diagnostic only.

Area	Key metric	Value	Caution
Rannacher gate	Recommendation	KEEP_ BASELINE_FOR_ NOW	Rannacher remains comparison tool.
Grid sensitivity	Max selected price diff vs reference	0.01111111111111	Nearest-node and three-grid limitation.
Domain sensitivity	Max selected price diff vs reference	0	Comparable ΔS design only.

7 Price Validation Evidence

The price evidence has three layers. First, Ticket 01 validates the closed-form European utilities and payoff functions. Second, Ticket 03 validates European CN against the closed-form formulas before adding American projection. Third, Tickets 05 through 07 validate American put, no-dividend call, and dividend-paying call behavior.

For beginner readers, this order matters. If European CN had failed, debugging PSOR would be premature. If the no-dividend American call had produced large artificial early-exercise value, the obstacle machinery would be suspicious. Instead, the evidence is internally consistent: the European solver is close to closed form, the American put shows early-exercise value, the no-dividend call remains close to European value, and the dividend-paying call shows positive American-over-European value.

8 American Obstacle and LCP Evidence

The American option value must satisfy $U \geq \Phi$, where Φ is the payoff. Ticket 04 implemented projection through PSOR, and Tickets 05, 07, 08, and 11 checked that projection behavior.

The strongest structural evidence comes from the LCP diagnostics. For the representative American put and dividend-paying call, the maximum obstacle violation is 0. The maximum equation violation reported by Ticket 08 is approximately 2.14×10^{-9} , and the maximum complementarity product is approximately 6.65×10^{-11} . These are near solver tolerance and support the claim that the implemented CN/PSOR core respects the discrete obstacle problem in the tested cases.

This is not a mathematical convergence proof. It is implementation evidence: the computed grids satisfy the intended discrete conditions closely for the validation ladder cases.

9 Boundary Extraction Evidence

Ticket 09 is the first ticket where boundary extraction is allowed. It uses the continuation premium

$$P_{\text{cont}}(S, \tau) = U(S, \tau) - \Phi(S)$$

and extracts threshold-based approximate boundaries. The American put and dividend-paying American call each produce 120 found positive- τ boundary rows in the representative $M = N = 120$ cases. The no-dividend call control produces `NO_BOUNDARY_FOUND`, which is important because a boundary extractor that always finds something would be dangerous.

The supported claim is modest: boundary extraction is useful as a diagnostic and research object. The

unsupported claim would be exact analytical free-boundary accuracy. Threshold choice, grid spacing, interpolation, and near-maturity behavior remain limitations.

10 Greek Diagnostic Evidence

Ticket 10 computes finite-difference Delta and Gamma. Delta and Gamma are useful because risk surfaces will eventually need them, but they are also fragile. Gamma is a second derivative, so it amplifies payoff kinks, free-boundary transitions, and grid noise.

The Ticket 10 summary reports full-grid $\max |\Gamma| = 30$ for the representative put and dividend call. After excluding maturity rows, boundary-near nodes, kink-near nodes, and endpoints, the strict-mask $\max |\Gamma|$ is much smaller, with maximum reported value 3.18001886582 in the Ticket 12 evidence table. Ticket 11 shows that Gamma remains grid-sensitive: finer grids expose sharper curvature.

Therefore, Greeks are validated as diagnostics, not as production-quality labels. They should not yet be exported as machine-learning targets.

11 Rannacher Smoothing Gate Evidence

Ticket 10A tested a named Rannacher-smoothed variant with two backward-Euler half-substeps over the first ordinary CN time interval. It did not silently replace the Ticket 04 baseline solver.

The comparison found very small baseline-vs-Rannacher price differences, with maximum absolute difference $1.72323866424 \times 10^{-6}$. Boundary shifts were small in the Ticket 10A report, but strict-mask Gamma did not improve enough to justify adoption. The gate recommendation was `KEEP_BASELINE_FOR_NOW`.

Ticket 12 respects that decision. The solver-validation synthesis uses the Ticket 04 baseline as the main validated solver.

12 Grid and Domain Sensitivity Evidence

Ticket 11 checks whether the representative results remain stable when grid resolution and domain cutoff change. Grid sensitivity used $S_{\max} = 4$ and $M = N = 80, 120, 180$. Domain sensitivity used $S_{\max} = 4, 5, 6$ with $M = 120, 150, 180$, keeping ΔS comparable and $N = 120$.

The selected-price differences versus the finest grid are small but not zero. The maximum selected price difference in grid sensitivity is 0.011111111111. Domain sensitivity is flatter in the selected target region under comparable ΔS , with maximum selected price difference 0. Boundary shifts exist under grid refinement, and Gamma remains grid-sensitive, so the evidence should be read as a preliminary sensitivity gate rather than a formal convergence proof.

13 Artifact Audit Summary

The Ticket 12 audit checked 113 expected artifacts and found 0 missing. It did not create replacement evidence for absent files.

Artifact group	Checked	Missing	Interpretation
Project documents	5	0	Source-basis documents are present.
Ticket reports	24	0	Ticket 01 through 11 L ^A T _E X and PDF reports are present.
Source files	10	0	Solver and diagnostic source modules through Ticket 11 are present.
Tests	12	0	Test files through Ticket 11 are present.
Experiments	9	0	Validation/audit scripts through Ticket 11 are present.
Result tables	24	0	Prior CSV result artifacts are present.
Figures	29	0	Prior diagnostic figures are present.
Total	113	0	ALL_REQUIRED_PRESENT.

14 Gate Decision

The Ticket 12 gate decision is:

PASS_SOLVER_VALIDATION_WITH_LIMITATIONS

Item	Decision record
Decision	PASS_SOLVER_VALIDATION_WITH_LIMITATIONS.
Evidence status	COMPLETE_WITH_LIMITATIONS.
Artifact status	ALL_REQUIRED_PRESENT.
Recommended next stage	Formal reference integration and human review before downstream risk-surface work.
Blocked until later	Stress maps, datasets, neural surrogates, production Greek labels, and final bibliography.
Main reason	The validation ladder is internally consistent and no Ticket 10A or Ticket 11 evidence contradicts continued baseline use.

This is intentionally not the stronger decision that would send the project directly into the next research stage. The project can move to reference integration and human review, but the phrase “with limitations” should remain visible because later research stages depend on preserving those cautions.

15 Claims Supported by Current Evidence

Supported claim	Evidence basis
European CN mode matches closed-form prices to finite-difference accuracy in the tested target region.	Ticket 03 closed-form error CSV.
The American CN/PSOR solver respects the payoff obstacle in tested representative cases.	Tickets 05, 07, 08, and 11 obstacle diagnostics.
The no-dividend American call behavior is consistent with the theorem that early exercise is not optimal.	Ticket 06 American-European comparison and Ticket 09 no-boundary control.
The dividend-paying American call case can show positive American-over-European value.	Ticket 07 positive-dividend validation.
The LCP/complementarity residuals are small in representative cases.	Ticket 08 diagnostics and Ticket 11 diagnostic sensitivity.
Continuation-premium boundary extraction works as a threshold-based diagnostic.	Ticket 09 boundary extraction and Ticket 11 boundary sensitivity.
Finite-difference Greeks can be computed and diagnosed with masks.	Ticket 10 Greek diagnostics and Ticket 11 Gamma sensitivity.
The baseline solver remains the default after the Rannacher comparison gate.	Ticket 10A recommendation KEEP_BASELINE_FOR_NOW.

16 Claims Not Yet Supported

Unsupported or premature claim	Why it is not yet supported
The solver has a full convergence proof across all parameters.	Ticket 11 uses three grid levels and selected domain cases only.
Extracted boundaries are exact analytical free boundaries.	Ticket 09 extraction is threshold-based and grid-dependent.
Delta and Gamma are production-quality risk measures.	Ticket 10 reports Gamma fragility near kinks, maturity, and boundary bands.
Rannacher smoothing should replace the baseline solver.	Ticket 10A explicitly keeps the baseline for now.
Stress maps are ready.	Stress maps require broader parameter validation and human review.
Dataset labels are reliable enough for neural training.	Greek and boundary labels remain diagnostic and need reference integration.
The final literature review is complete.	Formal reference integration is deferred.

17 Implementation Pitfalls and Debugging Notes

Actual implementation notes for Ticket 12:

Iteration	Symptom	Fix	Lesson
Initial red run	The focused Ticket 12 tests failed because <code>experiments/10_solver_validation_synthesis.py</code> did not exist.	Implemented the synthesis/audit script.	A red run should fail for the intended missing capability.
Audit label check	Simple report enumeration would have mislabeled Ticket 11 after Ticket 10A.	Derived ticket labels from artifact filenames.	Special ticket numbers such as 10A need explicit handling.

Likely synthesis pitfalls for future work:

- overclaiming solver reliability from representative cases;
- treating threshold-based boundaries as exact free boundaries;
- treating finite-difference Greeks as production risk labels;
- treating sensitivity outputs as datasets;
- ignoring missing artifacts;
- confusing a validation gate with completion of the whole project.

18 Limitations

Limitation	Meaning
Representative cases	The ladder focuses on normalized $K = 1$ cases and selected put/call regimes.
No formal convergence proof	Grid/domain sensitivity is preliminary and empirical.
Threshold-based boundaries	Boundary curves are diagnostics, not exact analytical objects.
Diagnostic Greeks	Delta and Gamma are finite-difference diagnostics with masks and cautions.
No stress maps	Ticket 12 does not generate regime maps or broad parameter grids.
No datasets or neural labels	Machine-learning label generation remains blocked.
No final reference integration	Bibliographic polishing is deferred to a later pass.

19 Lessons Learned / Study Notes

For greenhand researchers, the main lesson is that solver validation is a ladder, not one big test. A low price error is not enough by itself. American option solvers also need obstacle checks, complementarity checks, financial theorem checks, boundary diagnostics, Greek cautions, and sensitivity checks.

A second lesson is that every diagnostic has a scope. Boundary extraction from $U - \Phi$ is useful, but threshold-based. Gamma is useful, but fragile. Rannacher smoothing is useful to compare, but should

not be adopted just because it sounds more advanced. Grid sensitivity is useful, but three grid levels are not a proof.

A third lesson is that a conservative gate can still be progress. `PASS_SOLVER_VALIDATION_WITH_LIMITATIONS` is meaningful because it tells the project what can be trusted, what must remain cautious, and what work must not start yet.

20 Recommended Next Steps

Next step	Recommendation
Reference integration	Review the 12 ticket reports against verified classical American option, finite-difference, LCP/PSOR, Rannacher, and scientific machine-learning references.
Human review gate	Have the research team inspect the validation evidence and limitations before downstream work.
Broader validation plan	Decide whether additional grids, maturities, volatilities, rates, and dividend yields are needed before stress maps.
Boundary and Greek policy	Keep threshold and mask metadata attached to any future boundary or Greek output.
Downstream work	Do not create stress maps, datasets, neural labels, or surrogate models until the review gate accepts the limitations.

21 Reproducibility Record

Item	Record
Audit script	<code>experiments/10_solver_validation_synthesis.py</code> .
Tests	<code>tests/test_solver_validation_synthesis.py</code> .
Evidence CSV	<code>results/01_solver_validation/tables/ticket_12_solver_validation_evidence_summary.csv</code> .
Gate CSV	<code>results/01_solver_validation/tables/ticket_12_validation_gate_decision.csv</code> .
Artifact audit CSV	<code>results/01_solver_validation/tables/ticket_12_artifact_audit.csv</code> .
Focused Ticket 12 tests	Ran 9 tests in 0.030s OK.
Full test suite observed before report writing	Ran 157 tests in 4.282s OK.
Report	This L ^A T _E X source and compiled PDF.